

THRIVEXLABS


THRIVEXLABS FIELD REPORT

Reply Rates Lie

A Field Report on Trust-Led Outbound for B2B Teams, and
What to Measure Instead

Mayank Mishra

Founder, ThriveXLabs



revops.eden.so

Executive Summary

Most B2B outbound is judged by the wrong number.

This report comes from running LinkedIn outbound at real volume during a 0-to-1 growth sprint for an expert marketplace, a sprint that took the platform to 12.5k users in three months. A five-person team sent connection requests with personal notes (**30–35% accepted**) and **11,466 cold DMs**, and **46% of those DMs got a reply**. By any usual standard, that's a success.

It wasn't, and the reason is the most important lesson in B2B outbound: **the constraint was never getting replies. It was turning replies into real bookings**. Volume was solved. Conversion wasn't.

11,466

cold LinkedIn DMs sent by a five-person team

46%

reply rate, which looked like success on the dashboard (30–35% of connection requests accepted)

After the reply

where the real bottleneck was: replies turning into bookings

What we learned, in brief:

1. **Reply rate is a vanity metric when it's reported alone.** It measures attention, not readiness. Teams that optimize for it tend to push harder, earlier, and it shows up as ghosting after the call request.
2. **Targeting is the first act of reading the buyer.** A company that just raised money is at a different decision point from one already posting freelance roles. Same message, completely different buyer state.
3. **Segment underperformance is rarely a copy problem.** Our worst-performing niche didn't need better words. It needed a different reading of who those buyers were and what they'd already been sent.
4. **Small, specific proof beats big claims.** One of the cleanest conversions we saw came from a single LinkedIn post with a booking link, with no sequence at all.
5. **Measurement discipline is a competitive advantage.** Blended numbers hide which motion is actually working.

The rest of this report explains each finding, the operating model we built from it (**trust-led outbound**), and ten questions to ask any team, internal or external, that runs outbound for you.

Mayank Mishra, Founder, ThriveXLabs

Reply Rates Lie

A reply rate tells you one thing: how many people found your message worth answering. It doesn't tell you how many of them had a problem you solve, were in a position to act, or trusted you enough to take the next step.

At 46%, our reply rate was far above any benchmark people usually cite. Our own internal assessment was blunt: **the bottleneck was conversion of replies into real bookings, not more sends**. More volume would only have produced more of the same gap.

When we read the threads instead of the dashboard, a pattern showed up again and again:

Prospect: Interesting, what do you do exactly?

Typical reply: [offer description] + [proof point] + [link] + "Would you be open to a quick call this week?"

The prospect asked one question. They got answers to five they hadn't asked, plus a request for thirty minutes of their week. Then they went quiet, and the follow-up sequence fired on a timer, as if nothing had happened.

The diagnosis: the prospect was at an early trust question ("*Is this relevant to me?*"). The reply jumped to a late one ("*Is this worth your time and money?*"). That's **premature pressure**: answering a question the buyer hasn't reached yet. It's the single most common way outbound burns the replies it worked hard to earn.

What this means for your team

- **Report reply rate next to what happens after it:** substantive conversations, qualified opportunities, evaluations and outcomes, each with a stated denominator.
- **Track refusals separately.** A reply rate that counts "not interested" as a reply flatters itself.
- **Watch what happens after the call request.** If threads go quiet right after the ask, the ask is premature. That isn't the prospects' fault.

Targeting Is the First Act of Reading the Buyer

We tested two ways of finding companies likely to need project-based growth help: **hiring signals** (companies posting freelance or contractor roles on job boards) and **funding signals** (companies that had recently raised a seed or Series A round).

Job-board targeting looked logical: these companies clearly need help. In practice it was structurally weak:

- **By the time a company posts a freelance role, it has already chosen how to solve the problem.** It has selected against alternatives. You're reaching a buyer who has already decided.
- **The data was thin.** Job-board results gave us company name, posting link and location, with no contact data. That meant a separate enrichment pass, which erased any speed advantage.

Funding announcements were a **pre-decision signal**. A company that has just raised has fresh budget and usually hasn't locked in how it will solve its next growth problem. That's the window to be in front of them.

The principle

The same message can be relevant to one buyer and irrelevant to another, depending on **where they are in their decision**, not on the wording. Choosing *who* and *when* is choosing the buyer's trust state before you've sent a word.

Qualification-note discipline

Every target on a list we built needed a one-line "**why this company**" note tied to the ideal customer profile. For example: *"Raised a seed round recently; no in-house marketer on LinkedIn; B2B SaaS."* A list without those notes decays into spray-and-pray, and spray-and-pray is exactly what produces high reply rates on low-readiness buyers.

Segments Aren't a Copy Problem, and Small Proof Beats Big Claims

Finding 3: underperformance usually isn't about wording

We ran outbound across three niches. One of them, no-code/AI automation, had the worst reply rate of the three by a clear margin.

The reflex was to rewrite the copy. The better reading: this audience was more technical, more skeptical of "growth" framing, and already flooded with AI-tool cold outreach. They were standing at a different trust question. They needed to be convinced the sender was **real and relevant** before anything else could land. Changing adjectives doesn't fix that. Changing the hook, the proof and sometimes the targeting does.

For your team: when one segment underperforms, ask *"What is this audience already skeptical of, and what have they already been sent?"* before you ask *"How do we reword this?"* Sometimes the right conclusion is that the segment is wrong.

We saw a related pattern in our supply-side sourcing: accounts with very large followings (20K+) were often **already monetizing elsewhere** and converted poorly. The biggest names aren't automatically the best prospects. Sometimes the signal that makes them look attractive is the same signal that means they've already chosen.

Finding 4: specific, small proof travels furthest

One of the cleanest conversion stories we saw involved no outbound at all. An expert added a booking link to a single LinkedIn post and received a booking.

It's a small story, and that's exactly why it's persuasive. It's specific, verifiable and believable. It also shows how the system really works: **content created the memory, the profile confirmed it, and the buyer moved on their own.** No sequence required.

For your team: your best proof is rarely your biggest claim. It's the most specific, true example that matches the buyer's situation. Inflated claims ("we generate 10–20 qualified leads a month") fail the buyer's value test the moment they meet your profile.

Measure Like You Want the Truth

Outbound dashboards are easy to make impressive and hard to make honest. We held ourselves to a few rules that I now recommend to every B2B team:

- 1. **Never blend what's different.** Keep different kinds of outcomes in separate columns: real versus subsidized, outbound versus inbound, one offer type versus another. A single blended number hides which motion is doing the work.
- 2. **Name the denominator.** "Reply rate" means unique people who replied ÷ unique people messaged, in a defined cohort and period. Follow-ups aren't new contacts.
- 3. **Show counts next to percentages.** A 50% win rate off two proposals is one data point, not a trend.
- 4. **Say "not yet measurable" when it isn't.** Substituting a proxy without labeling it is how teams talk themselves into the wrong strategy.
- 5. **Read threads before rewriting anything.** The dashboard tells you *where* it's leaking. Only the conversations tell you *why*.


Where the leak usually is

Symptom	Where to look first
Few accept connection requests	Fit, timing signal, profile clarity
Accepted, but no replies	Is the opener relevant and easy to answer?
Lots of polite replies, few substantive ones	Targeting and the substance of the question
Substantive, but rarely qualified	Priority, resources, a real buying situation
Qualified, but stalls before a call or evaluation	The next step you proposed, and what it asks of the buyer
Calls held, proposals lost	Proof, scope, price, alternatives, decision process

The first message's wording is often not the constraint.

Trust-Led Outbound

Everything above leads to one operating model. Outbound works when it is designed around the buyer's trust state at every step, not around the seller's sequence.



1. Signal-based targeting

Pre-decision signals (such as a recent funding round, a new leadership hire in the function, a public discussion of the exact problem) over post-decision signals (a job posting that shows the path is already chosen). Every target carries a written reason.

2. Openers that pass the first two trust questions

Buyers ask *"Is this real?"* and *"Is this relevant to me?"* before anything else. Openers lead with a specific, verifiable observation, state plainly what the sender does, and ask one question that the buyer can honestly answer "no." Connection requests don't pitch.

3. Trust-led conversations

This is where most outbound breaks and where the method in *The Non-Salesy DM* comes in. Every reply is read for:

- **State:** how the buyer is engaging (curious, skeptical, exploring, ready...)
- **Gate:** which of six trust questions they're answering: Reality, Relevance, Fit, Safety, Value, or Identity/Commitment
- **Risk:** what they're protecting (time, budget, reputation, a past bad experience)
- **The missing signal,** and the **smallest move** that supplies it

One message, one job. No stacked pitches. No call requests to people who haven't decided you're relevant. Sometimes the right move is to send nothing.

4. Clear offers when the buyer is ready

The opposite failure matters just as much: **hiding**. When a buyer asks how you'd approach their problem, they get a clear recap, a scoped offer and a price. They don't get another "helpful resource."

5. Measurement that doesn't flatter

Leading indicators (substantive replies, volunteered context, calls that don't go quiet afterward) are reported next to lagging ones (qualified pipeline, proposals, revenue). Denominators are stated. Different motions are never blended.

Ten Questions to Ask Anyone Who Runs Outbound for You

Whether you're hiring an agency, building an internal team or reviewing your own setup:

1. **What do you report besides reply rate?** Look for substantive conversations, qualified opportunities and outcomes, each with a denominator.
2. **How do you choose who to contact, and why now?** Listen for pre-decision signals and written qualification notes, not "job title plus company size."
3. **What does a connection request say?** If the answer involves a pitch, the program leads with pressure.
4. **What happens when a prospect replies "interesting"?** Listen for judgment. A pre-scheduled next message is a red flag.
5. **When do you ask for a call, and what happens to threads after you do?** Ask to see how many go quiet right after the call request.
6. **How do you handle "not now" and "no"?** Listen for recorded triggers and hard stops, not "we follow up until they reply."
7. **Show me three real threads**, including one that didn't convert. Judge the conversations, not the dashboard.
8. **What proof do you use, and can every claim be verified?** Inflated claims fail buyers' value test.
9. **How do you keep different outcome types separate in reporting?** Blended numbers hide what's working.
10. **What would make you tell me outbound isn't the constraint?** A good operator knows when the real problem is the offer, the ideal customer profile or the profile page.

Work With ThriveXLabs


At **ThriveXLabs**, I build and run outbound-led sales systems for B2B businesses using the trust-led model in this report: signal-based targeting, conversations that read the buyer instead of running a script, clear offers when buyers are ready, and measurement that tells you the truth.

Two ways to start

Run it yourself. *The Non-Salesy DM* teaches the conversation method behind this report. There's a free Field Guide, and a Complete Edition with a tracker, an AI thread debugger and a 30-day sprint.

Have it run for you. Talk to ThriveXLabs about building a trust-led outbound system for your team.

Everything is at revops.eden.so



Scan to visit
revops.eden.so

The figures in this report (11,466 cold LinkedIn DMs at a 46% reply rate, a 30–35% connection-acceptance rate on requests with notes, and a platform that grew to 12.5k users in three months) come from the author's own operating experience running outbound for one business during one period. They are not benchmarks or promises. Teaching examples are illustrative. LinkedIn is a trademark of its owner; this report is independent. © 2026 Mayank Mishra · ThriveXLabs.